

Integration testing for Jakarta EE

Rudy De Busscher

Version 2.2.0, 10/04/2023

Table of Contents

Release notes	1
2.2.0	1
1.2.0	1
1.1.0	1
1.0.0	1
Introduction	2
Supported Jakarta runtimes	3
Configuration	4
Creating a test	5
Access application endpoints	6
Base Docker images	7
Defining version number	8
Custom Docker Build scripts (v1.1.0)	9
Payara Micro	9
OpenLiberty	9
Wildfly	9

Release notes

2.2.0

1. Jakarta EE 10 based of version 1.2.0
2. Dropped Support for Glassfish (as there is no Docker image available in public repository).
3. OpenLiberty support is not available as there is no Jakarta EE 10 version available. Should be available soon.

1.2.0

1. WAR file can be searched in other directories than *target*, include some path outside of the project with `applicationLocation` member of the test annotation.
2. Adapter functionality for Docker Build scripts
3. Support for testing applications that access a database.
4. Update dependencies to the latest stable version.

1.1.0

1. Added support for custom Docker build scripts.
2. Easier for using additional containers that are used by the application under test (microservices set-up)
3. Using *WireMock* for providing responses of remote services.
4. Update version of Jakarta runtimes.

1.0.0

1. Initial version with support for Payara Micro, Open Liberty, Wildfly
2. Glassfish support is limited.

Introduction

Integration testing can be difficult as you need to run your application as it will run in production. This means that you need to set up and configure all dependencies but provide each of them with a predefined set of values to have repeatable and predictable tests.

The use of the [Testcontainers project](#) can help you in setting up and maintaining the test environments that you need.

The Jakarta Integration Tet framework helps you to run your application on the application container within Docker using the Testcontainers project. It helps you deploy the application with the runtime of your choice and provides you easy access to the endpoints of your application. But with a simple parameter, you can also remotely debug your application to get a fast solution to the issues that might be detected during the tests.

Supported Jakarta runtimes

This integration testing framework is designed to work with Jakarta Runtimes, Jakarta EE 8 as this is still the version that is used by the bulk of the applications that are developed today, but actually it is not limited to this version (see further on). The runtimes that are supported are:

- Payara Micro - [website](#)
- OpenLiberty - [website](#)
- Wildfly - [website](#)

Configuration

To get started, add this dependency to your project

```
<dependency>
  <groupId>be.atbash.test</groupId>
  <artifactId>integration-testing</artifactId>
  <version>2.2.0</version>
  <scope>test</scope>
</dependency>
```

It uses JUnit5 and also brings the Testcontainers dependency into your project test scope. It also depends on Apache CXF and Jackson but more on that when I talk about calling the endpoints within your application.

On your machine, you also need a valid Docker client (although it also works with Docker Machine) that accesses a docker environment. Look at the Testcontainers requirement for Docker if you want more information.

Creating a test

When you want to write a test that deploys your application on a certain runtime, you need to have the following test class.

```
@ContainerIntegrationTest(runtime = SupportedRuntime.PAYARA_MICRO)
public class HelloPayaraIT extends AbstractContainerIntegrationTest {
    // Should be public and not the JUnit 5 preferred package scope.
}
```

The `AbstractContainerIntegrationTest` class keeps the Container reference that is running your application and also registers an extension that will help you trace problems (see further on)

The class also requires a specific annotation that triggers the JUnit 5 extension but also has some members to define, for example, the Jakarta runtime on which your application will run. For an explanation of all members, see further on

As you can see from the comment in the snippet, the test class should have a scope `public`. The JUnit 5 extension makes use of reflection and to avoid issues on newer JDK versions, the accessibility of fields is not changed (no call of `field.setAccessible(true)`).

Access application endpoints

When you want to perform integration testing on your application, you typically call the endpoints and see if you get the expected results back. Calling endpoints or URLs can be done in several ways within Java, even with the classes from the JVM itself. But most methods are rather cumbersome and sensitive to configuration values like host, ports, paths, etc. And these vary which each run using Testcontainers, so a robust way of calling your endpoints was chosen using the MicroProfile rest Client functionality.

But as a developer writing the test, you don't need to know all the details to make use of it nor does the runtime that runs your application must support MicroProfile Rest Client since it is only used on the test side.

For each *controller* or the set of endpoints you want to test, you should create an interface with the JAX-RS annotations that define the endpoint.

```
@Path("/product")
@Consumes(MediaType.APPLICATION_JSON)
@Produces(MediaType.APPLICATION_JSON)
public interface JsonService {

    @GET
    List<Product> getProducts();

    @POST
    void addProduct(Product product);

}
```

With the help of the Apache CXF and Jackson frameworks, the JUnit 5 extension will generate a proxy from this interface that is capable of calling your application endpoint. You only need to *inject* it into your test class. The JUnit5 extensions use the RestClient annotation as a marker for this purpose.

```
@RestClient
public JsonService jsonService;
```

(also here the field must be public to avoid calls to `setAccessible`)

Calling your application endpoints in the test becomes calling Java methods. And the host, port, and root part of your URL is taken care of by the JUnit 5 extension. You only need to configure the rest of the path with the `@Path` annotation on the interface.

Base Docker images

As you have seen in the example earlier in this document, you can define the runtime that runs your application using the `runtime` member of the `@ContainerIntegrationTest` annotation.

But you also can define the runtime using a java system property. When you do not specify the runtime with the annotation, the property `be.atbash.test.runtime` is used to determine the runtime. The value should, case-insensitive, match the enum name of `SupportedRuntime`. This allows you to run your application on different runtimes if you are developing a framework or library for example.

The value of the `SupportedRuntime` determines the base Docker image from which the JUnit 5 extension creates an Image to perform the test. These are the default values of those Docker images.

- Payara Micro : `payara/micro:6.2023.2`
- OpenLiberty : Not available yet
- WildFly : `quay.io/wildfly/wildfly:27.0.0.Final-jdk11`

You can use a different base docker image by defining a System Property.

Defining version number

When you specify the System property `be.atbash.test.runtime.version`, it is used to determine the base Docker image.

You can specify just the tagname to select the same Docker image but another version (like `6.2023.2-jdk17` for the Java 17 version of the Payara Micro image). When this value contains a `:` or `/`, it will be used as the value for the *FROM* command in the Docker file. This allows you to use your own Docker image for testing your application.

This is ignored when using a custom Docker build script.

Custom Docker Build scripts (v1.1.0)

Instead of using the default images for the supported runtimes as mentioned earlier, you can also define a custom build script that will be used.

To indicate the directory that contains the Docker build artifacts, use the annotation `@CustomBuildFile` annotation on your test class.

```
@ContainerIntegrationTest(runtime = SupportedRuntime.PAYARA_MICRO)
@CustomBuildFile(location = "custom/payara")
public class CustomPayaraIT extends AbstractContainerIntegrationTest {
```

You still need to indicate the runtime as some runtime specific actions are performed (this can also be through the System property) The location is relative to the `<project-root>/src/docker` directory.

The directory can contain a file called `Dockerfile` that will be used as build for the image.

The directory itself and all subdirectories are also included in the build tar, so it can contain additional files referenced within the Docker build file.

If no file called `Dockerfile` is found, the default one is used. In this case only the additional files are included into the Docker Image build but in most cases this will not be useful.

The following statements are added to this `Dockerfile` depending on the supported runtime.

Payara Micro

```
CMD ["--deploy", "/opt/payara/deployments/test.war", "--noCluster", "--contextRoot",
"/"]
ADD test.war /opt/payara/deployments
```

OpenLiberty

```
ADD test.war /config/apps
```

Wildfly

```
ADD test.war /opt/jboss/wildfly/standalone/deployments
```

== Adapter for Docker Build scripts (v1.2.0)

An alternative for providing a custom build script, is to define an adapter that can modify the contents of the build script just before it is sent to the Docker engine.

The adapter receives the default build script or the custom build script when the annotation is defined, and can adapt it. The return value is used as build script. And thus it is possible that a complete new content is used. This is up to the developer to decide what is needed.

An adapter must implement the interface `'be.atbash.testing.integration.container.image.DockerImageAdapter'` which has 2 methods.

[source]

```
public interface DockerImageAdapter { String adapt(String dockerFileContent, TestContext testContext);
```

```
    default SupportedRuntime supportedRuntime() {  
        return SupportedRuntime.DEFAULT;  
    }  
}
```

Implementations of this interface needs to be 'registered' through the Service Loader mechanism of Java.

The return value of the method `'supportedRuntime()'` method determines if the adapter is used for any supported runtime or only for a specific one.

There is the possibility to define multiple adapters and in that case the `@Priority` (from Jakarta EE 10) should be used to determine the order the adapters are applied. Lower values are applied first and when no annotation is specified, a value of 100 is assumed.

The instance `'TestContext'` provides the Adapter the information that is specified on the test annotations and the derived 'metadata'. The method parameter should be used as.

[source]

```
ContainerAdapterMetaData metaData =  
testContext.getInstance(ContainerAdapterMetaData.class);
```

When the database integration test is used, the `'ContainerAdapterMetaData'` will also be available within this context.

You can also provide additional Environment variables that will be made available on the container running your application. To do this, take the `AdditionalEnvParameters` instance from the `'TestContext'` and add the values.

[source]

```
AdditionalEnvParameters parameters =  
testContext.getInstance(AdditionalEnvParameters.class);  
parameters.add(key, value);
```

If that instance is not available yet on the context, you can add it.

== Using WireMock for fake remote responses (v1.1.0)

Many times your application calls other services to have all data to respond to the user request. These data can be provided in your test by putting a `_WireMock_` server as the server that provides you results for these remote calls.

Using `_WireMock_` is simplified within the Integration Testing framework as it has a specific class and methods to define the behaviour of the remote endpoints.

You can use a `_WireMock_` instance in your test by adding the following code snippet for defining the container.

```
@Container  
public static final WireMockContainer wireMockContainer =  
WireMockContainer.forHost("wire");
```

And defining the response for a certain URL call be done by using the following statements in your tests

```
MappingBuilder mappingBuilder = new MappingBuilder()  
    .forURL("/path")  
    .withBody(foo);
```

```
String mappingId = wireMockContainer.configureResponse(mappingBuilder);
```

Your method under test is then assumed to call a URL `'http://wire:8080/path'` as part of its logic. You can use the MicroProfile config and Rest client functionality, supported by Payara Micro, OpenLiberty and WildFly, to call and configure this endpoint.

The `'foo'` instance is a POJO that needs to be returned as the JSON response.

The `'MappingBuilder'` class has several methods to define the response of `_WireMock_`.

- `_.withBody()` with a String defines the body content and sets the `_Content-Type_` to `_text/plain_`. When the parameter is any other object type, it is converted to JSON and the `_Content-Type_` to `_application/json_`.
- `_.forURL()` define the path of the URL that will provide the response.
- `_.withStatus()` can be used to change the default return status 200.
- `_.withMethod()` defines the HTTP method for the request that will be supported. By default this is `_GET_`.
- `_.withContentType()` can be used to define a specific `_Content-Type_`. You should use this `_.withContentType()` only after setting the body as that method already sets a specific `_Content-Type_`.

After a test method is executed, the mapping configuration of the `_WireMock_` server is reset so that you can test different scenarios within one test class.

With the `_mappingId_` value that is returned by the `'configureResponse'` method, you can retrieve information about the request received by `_WireMock_`. The method `'WireMockContainer.getRequestInfo()'` method returns `'null'` when no request is received for that mapping, or some info if `_WireMock_` received a request for that mapping.

== Using Database (v1.2.0)

When your application or micro-service makes use of a database through JPA, you can configure the integration test to have a database started before the test with a known state. The settings for the database connections are placed as environment variables of the test application container.

The feature can be activated by using the `'@DatabaseContainerIntegrationTest'` annotation on a test class that extends `'AbstractDatabaseContainerIntegrationTest'` class.

Also, a database container must be added to the project dependencies. The database is discovered automatically and functionality is adapted according to the discovered database.

For example, the following dependency

[source]

```
<dependency>
  <groupId>org.testcontainers</groupId>
  <artifactId>mysql</artifactId>
  <version>1.17.6</version>
  <scope>test</scope>
</dependency>
```

triggers the support for MySQL. Other supported databases are PostgreSQL, MariaDB and Oracle XE.

The JDBC Driver that will be installed together with your application on the runtime, is determined from the selected database container and must be added to the project as test dependency.

The '@DatabaseContainerIntegrationTest' annotation allows for several configuration options. At a minimum, the 'ContainerIntegrationTest' must be specified, as with any Jakarta Integration test.

[source]

```
@DatabaseContainerIntegrationTest(                      containerIntegrationTest                      =
@ContainerIntegrationTest(runtime = SupportedRuntime.PAYARA_MICRO) )
```

Other options are

- _environmentParametersForDatabase_ defines the environment variables names that are used to transfer the JDBC URL, username and password for the database.
- _databaseScriptFiles_ defines the file names that are loaded from the class path to create the database tables and the Excel file that contains the data (records) that will be loaded for the test (through DBUnit).
- _databaseContainerImageName_ defines the Docker image name if you do not want to make use of the default image that comes with the dependency.
- _databaseContainerStartInParallel_ defines if the database starts in parallel with the other containers or separately as the first container.
- _jndiDatasourceName_ is used when using WildFly and must match the JNDI name you have specified within the _persistence.xml_ file.

Besides the start of the database container when the container running your application starts, it provides you also a way query the database during the test itself.

For that purpose, a 'IDatabaseConnection' from _DBUnit_ is accessible to test the changes that should be done by the application code during the test.

[source]

```
ITable companyTable = databaseConnection.createQueryTable("company", "SELECT id, name  
FROM Company");
```

The above statement retrieves all records in the `_Company_` table that can be used to verify if the call to an endpoint of the application created a new record as intended.

Depending on the runtime, a certain way of defining the datasource is assumed. They correspond with some best practices in a cloud environment so that configuration happens on the fly when the container starts.

=== Payara Micro

The datasource should be configured through `@DataSourceDefinition` and the url, user name and password configured through `MicroProfile` so that the values set on the environment are picked up and point to the started database container.

[source]

```
@DataSourceDefinition(name="java:global/jdbc/cloud-mysql",  
className="com.mysql.jdbc.jdbc2.optional.MysqlDataSource", url = "${MPCONFIG=ds_url}",  
user="${MPCONFIG=ds_username}", password="${MPCONFIG=ds_password}", maxPoolSize = 4,  
minPoolSize = 2)
```

The keys `_ds_url_`, `_ds_username_`, and `_ds_password_` can be changed by the `'environmentParametersForDatabase'` config option mentioned earlier.

=== OpenLiberty

For the OpenLiberty runtime, the data source is defined within `'server.xml'` (as you would do with any other project). The library fileset must point to `_config/lib_` as that is the location where Atbash Integration test code places the JDBC driver.

[source]

```
<dataSource jndiName="jdbc/cloud-mysql">  
  <jdbcDriver libraryRef="jdbcLib"/>  
  <properties server.env="true">  
    <url>${env.ds_url}</url>  
    <user>${env.ds_username}</user>  
    <password>${env.ds_password}</password>  
  </properties>  
  <connectionManager maxPoolSize="4" minPoolSize="2"/>  
</dataSource>
```



```
<library id="jdbcLib">
  <fileset dir="/config/lib" includes="*.jar"/>
</library>
```

The keys `_ds_url_`, `_ds_username_`, and `_ds_password_` can be changed by the `'environmentParametersForDatabase'` config option mentioned earlier.

=== WildFly

This is the runtime where the automatic support for database works the least. The datasource is deployed by creating a `_datasource XML_` file that is deployed alongside the application. But to have the variables replaced at startup, the docker image also need to install the Unix command `'envsubst'` which takes a bit of time.

Also, the JNDI name you specified in `th persistence.xml` must be defined through the `_jndiDatasourceName_` option mentioned earlier.

== Additional containers for a test (v1.1.0)

Additional containers can be started together with the container running your application under test.

All `'public static'` fields that are discovered in your test class having a `'@Container'` annotation and are of course a `Testcontainer`, are started together with the main container.

With version 1.1.0 there is an addition to the library to make it easier to add additional containers with applications when they are running on one of the supported runtimes.

[source, java]

`@ContainerIntegrationTest(runtime = SupportedRuntime.PAYARA_MICRO)` class `ApplicationTest` extends `AbstractContainerIntegrationTest` {

```
@Container
public static GenericContainer<?> remote = new
PayaraMicroContainer(DockerImageName.parse("test-remote:1.0"));
}
```

You can use the classes `'PayaraMicroContainer'`, `'OpenLibertyContainer'`, and `'WildflyContainer'` to run an additional image but already apply some logic to the container.

- The container shares the same network as the container running your application under test.
- The health check is automatically applied to the container and thus your test will not start until container reports itself as healthy.
- The variable name is added as alias host to the container so that the endpoints in this additional container can be accessed through a consistent, constant host name.

There is of course no need that the runtime of the container running your test is the same as the additional containers you start for your application.

== Define directory of WAR file (v1.2.0)

From version 1.2.0 onwards, you can define the directory where the WAR file is located. By default, it still looks i the target folder of your Maven project. But with the `'applicationLocation'` member of the annotation, you can define another location.

This can be another module of a multi-module Maven project or an absolute path.

[source, java]

```
@ContainerIntegrationTest(runtime = SupportedRuntime.PAYARA_MICRO, applicationLocation =  
"app") class ApplicationTest extends AbstractContainerIntegrationTest {  
  
}
```

The above example looks within the module directory `_app_` and the tests can be in then module `_tests_`.

project |-- app `-- tests

== Jakarta runtime specifics

This version of the framework is designed to work with any Jakarta EE 10 compatible runtime and the application is running with Java 11. For a few runtimes, there are some specific additional requirements or limitations.

=== OpenLiberty

The Docker images for OpenLiberty require that you supply a `_server.xml_` file to configure the process. The testing framework expects this file within the `_src/main/liberty/config_` directory (the standard location when using the liberty tooling.) For more information on this file, look at <https://openliberty.io/docs/latest/reference/config/server-configuration-overview.html> [OpenLiberty documentation] page and the examples in this repository also have a minimal example.

Important here is the element `'webApplication'` that makes sure the application under test is deployed on the root.

== Remote Debug

The testing framework supports remote debugging of your application. This makes it easier to research what is wrong with your code based on a failing test.

To activate it, set the `'Debug'` member of the `'@ContainerIntegrationTest'` annotation to true.

During the time that the framework code waits until the application is up and running, you can connect your Java debugger to port 5005. The start of the JVM is halted due to the `'suspend=y'` option that is passed to the JVM as part of the Debug configuration.

If you do not connect the debugger 'on time', the Testing framework reports the test as failed because the container did not start up correctly within 60 seconds of waiting time.

== Other features

Some additional features available with the testing framework

=== Volume Mapping

It is also possible to define a volume mapping between the host running the test and the container running the application. This is the easiest way when you need to send to or retrieve files from the container. The mapping can be defined within the `'@ContainerIntegrationTest'`

[source, java]

```
@ContainerIntegrationTest(volumeMapping = {"path/on/host", "/path/within/container"})
```

You can define 1 or multiple mappings by defining sets of 2, 4, 6, ... strings.

The first one is the directory on the host. It can be a relative path and is resolved against the current directory of the current process. It might also be an absolute path and the JVM logic is used to derive the absolute path for the value you specify (using `File.getAbsolutePath`).

The second string is the directory within the container and must always be absolute.

=== Live logging

It is possible to show the output of the runtime in the test output log. To have this info, specify it through the annotation. You should have already the logging for Testcontainers set up probably to have this working.

First, let us quickly recap the logging configuration of TestContainers. You can also read more on the

[https://www.testcontainers.org/supported_docker_environment/logging_config/\[Testcontainers documentation page\]](https://www.testcontainers.org/supported_docker_environment/logging_config/[Testcontainers documentation page]).

Make sure you add an SLF4J logging output dependency to your project, like `_Logback_`.

[source, xml]

```
<dependency>
  <groupId>ch.qos.logback</groupId>
  <artifactId>logback-classic</artifactId>
  <version>1.2.11</version>
  <scope>test</scope>
</dependency>
```

And have a proper configuration file for Logback appenders.

[source, java]

`@ContainerIntegrationTest(liveLogging = true)`

With the above definition, the container log will show within the output as defined in the Logback configuration at the moment the log entry is generated.

You can always access the log of the container from within your test code by using the statement.

[source, java]

```
String logs =  
AbstractContainerIntegrationTest.testContainer.AbstractContainerIntegrationTest.testCo  
ntainer.getLogs();
```

=== Container log when test fails

Without any additional configuration needed, the container log will be shown in certain cases of failure of your test. This will help you to determine what went wrong and how you can fix the problem. With the following types of failures, the log is shown.

- The test code throws a `'java.lang.AssertionError'` error.
- The MicroProfile Rest client code encounters a status 404 when calling an endpoint.
- The MicroProfile Rest client code detects an Internal Server Error within the container.

=== Jakarta EE 10 support

As indicated, the current version of the test framework runs runtime versions of Jakarta EE 8 compatible products on JDK 11.

But since the only `_connection_` between the test and the endpoints of your application within the container is HTTP based, there is no requirement on the application that can be tested.

When you define the version/tag name of the container that is started your application can make use of Jakarta EE 9.x, Jakarta EE 10, and run on any JDK that is supported by the runtime. So it is easy to use this framework with the upcoming Jakarta EE 10 release.